

Last Time

- Internet Application Security and Privacy
 - Network-layer security: VPN, IPSec
 - Transport-layer security and privacy: TLS / SSL

This time

[Stinson, Shmatikov-Boneh]

- Internet Application Security and Privacy
 - Transport-layer security and privacy: Tor
 - Application-layer security and privacy: **SSH**, **remailers**, PGP/gpg, OTR

Tor

- **Tor** is a successful privacy enhancing technology that works at the transport layer
 - Hundreds of thousands of users
- Normally, any TCP connection you make on the Internet automatically reveals your IP address
 - Why?
- Tor allows you to make TCP connections **without** revealing your IP address
- It's most commonly used for HTTP (web) connections

- Client connects to server, indicates it wants to speak TLS, and which **cipher suites** it knows
- Server sends its certificate to client, which contains:
 - Its host name
 - Its public key
 - Some other administrative information
 - A signature from a Certificate Authority
- Server also chooses which cipher suite to use
- Client sends symmetric encryption key K , *encrypted with server's public key*
- *Communication now proceeds using K and the chosen ciphersuite*


How Tor works

- Scattered around the Internet are about 1000 Tor **nodes**, also called **Onion Routers**
- Alice wants to connect to a web server without revealing her IP address




How Tor works

- Alice picks one of the Tor nodes (n1) and uses public-key cryptography to establish an encrypted communication channel to it (much like TLS)




How Tor works

- Alice tells n1 to contact a second node (n2), and establishes a new encrypted communication channel to n2, tunnelled within the previous one to n1




How Tor works

- Alice tells n2 to contact a third node (n3), and establishes a new encrypted communication channel to n3, tunnelled within the previous one to n2



How Tor works

- And so on, for as many steps as she likes (usually 3)
- Alice tells the last node (within the layers of tunnels) to connect to the website



Sending messages with Tor

- Alice now shares three symmetric keys:
 - $K1$ with $n1$
 - $K2$ with $n2$
 - $K3$ with $n3$
- When Alice wants to send a message M , she actually sends $E_{K1}(E_{K2}(E_{K3}(M)))$
- Node $n1$ uses $K1$ to decrypt the outer layer, and passes the result $E_{K2}(E_{K3}(M))$ to $n2$
- Node $n2$ uses $K2$ to decrypt the next layer, and passes the result $E_{K3}(M)$ to $n3$
- Node $n3$ uses $K3$ to decrypt the final layer, and sends M to the website

Replies in Tor

- When the website replies with message R , it will send it to node $n3$
 - Why?
- Node $n3$ will **encrypt** R with $K3$ and send $E_{K3}(R)$ to $n2$
- Node $n2$ will encrypt that with $K2$ and send $E_{K2}(E_{K3}(R))$ to $n1$
- Node $n1$ will encrypt that with $K1$ and send $E_{K1}(E_{K2}(E_{K3}(R)))$ to Alice
- Alice will use $K1$, $K2$, and $K3$ to decrypt the layers of the reply and recover R

Who knows what?

- Notice that node n1 knows that Alice is using Tor, and that her next node is n2, but does not know which website Alice is visiting
- Node n3 knows some Tor user (with previous node n2) is using a particular website, but doesn't know who
- The website itself only knows that it got a connection from Tor node n3
- **Note:** the connection between n3 and the website is **not encrypted**! If you want encryption as well as the benefits of Tor, you should use encryption **in addition** (HTTPS)

Anonymity vs. pseudonymity

- Tor provides for **anonymity** in TCP connections over the Internet, both **unlinkably** (long-term) and **linkably** (short-term)
- What does this mean?
 - There's no long-term identifier for a Tor user
 - If a web server gets a connection from Tor today, and another one tomorrow, it won't be able to tell whether those are from the same person
 - But two connections in quick succession from the same Tor node are more likely to in fact be from the same person

Application-layer security and privacy

- TLS can provide for encryption at the TCP socket level
 - “End-to-end” in the sense of a network connection
 - Is this good enough? Consider SMTPS (SMTP/email over TLS)
- Many applications would like true end-to-end security
- We'll look at three particular applications:
 - Remote login, email, instant messaging

Secure remote login (ssh)

- You're already familiar with this tool for securely logging in to a remote machine
- Usual usage (simplified):
 - Client connects to server
 - Server sends its public key
 - The client **should** verify that this is the correct key
 - Client picks a random **session key**, encrypts it with server's public key, sends to server
 - All communication from here on in is encrypted and MACd with the session key
 - Client authenticates to server
 - Server accepts authentication, login proceeds (under encryption and MAC)

Authentication with ssh

- There are two main ways to authenticate with ssh:
 - Send a password over the encrypted channel
 - The server needs to know (a hash of) your password
 - Sign a challenge with your private signature key
 - The server needs to know your public key
- Which is better? Why?

[<http://www.debian-administration.org/articles/530>]

Anonymity for email: remailers

- Tor allows you to anonymously communicate over the Internet in real time
 - What about (non-interactive) email?
 - This is actually an easier problem, and was implemented much earlier than Tor
- Anonymous remailers allow you to send email without revealing your own email address
 - Of course, it's hard to have a conversation that way
 - Pseudonymity is useful in the context of email

Recap

- Internet Application Security and Privacy
 - Transport-layer security and privacy: Tor
 - Application-layer security and privacy: SSH, remailers

Next time

- Internet Application Security and Privacy
 - Application-layer security and privacy:
remailers, PGP/gpg